



IMPLEMENTAREA CONCURENTEI IN LIMBAJE DE PROGRAMARE

Concurenta


Threaduri

Memorie Partajata

Ioana Leustean



[Part II. Concurrent Haskell](#)
Cap.7 & 8



`forkIO :: IO () -> IO ThreadId`

MVar
mutable variable

➤ Comunicarea folosind **MVar** se face in **monada IO**


data MVar a

newEmptyMVar :: IO (MVar a) -- m <- newEmptyMVar
-- m este o locatie goala

newMVar :: a -> IO (MVar a) -- m <- newMVar v
-- m este o locatie care contine valoarea v


takeMVar :: MVar a -> IO a -- v <- takeMVar m
-- intoarce in v valoarea din m
-- asteapta (blocheaza thread-ul) daca m este goala

putMVar :: MVar a -> a -> IO() -- putMVar m v
-- pune in m valoarea v
-- asteapta (blocheaza thread-ul) daca m este plina



- Sincronizare: serviciu de logare
modelarea unui canal de comunicare simplu folosind `MVar`


http://chimera.labs.oreilly.com/books/1230000000929/ch07.html#sec_conc-logger



Cerinte:

- serviciul de logare prelucreaza mesajele intr-un thread separat
- mesajele trebuie prelucrate in ordinea in care sunt logate
- cand programul se termina toate mesajele logate trebuie sa fie prelucrate

Exemplu: serviciu de logare – varianta1



```
logger :: Logger -> IO ()
logger (Logger m) = loop
  where
    loop = do
      msg <- takeMVar m
      putStrLn msg
      loop
```

```
data Logger = Logger MVar String
```

```
initLogger :: IO Logger
```

```
initLogger = do
```

```
    m <- newEmptyMVar
```

```
    let log = Logger m    -- log =
```

```
    forkIO (logger log)    -- creeaza
```

```
    return log
```

```
logger :: Logger -> IO() -- prelucreaza mesajele din Logger
```



Exemplu: serviciu de logare- varianta1




```
logMessage :: Logger -> String -> IO ()
logMessage (Logger m) s = putMVar m s

logMessTh :: Logger -> IO()
logMessTh l = do
    msg <- getLine
    if (msg == "bye")
    then return()
    else do
        logMessage log msg
        logMessTh log

main = do
    log <- initLogger
    logMessTh log
```

Thread-ul principal trimite mesajele



```
data Logger = Logger MVar String

initLogger :: IO Logger
initLogger = do
    m <- newEmptyMVar
    let log = Logger m
    forkIO (logger log)
    return log

logger :: Logger -> IO ()
logger (Logger m) = loop
    where
        loop = do
            msg <- takeMVar m
            putStrLn msg
            loop
```

Thread-ul logger le citește și le scrie




Exemplu: serviciu de logare- varianta1



```
logMessage :: Logger -> String -> IO ()
logMessage (Logger m) s = putMVar m s

logMessTh :: Logger -> IO()
logMessTh log = do
    msg <- getLine
    if (msg == "bye")
    then return()
    else do
        logMessage log msg
        logMessTh log

main = do
    log <- initLogger
    logMessTh log
```



```
data Logger = Logger MVar String

initLogger :: IO Logger
initLogger = do
    m <- newEmptyMVar
    let log = Logger m
    forkIO (logger log)
    return log

logger :: Logger -> IO ()
logger (Logger m) = loop
    where
        loop = do
            msg <- takeMVar m
            putStrLn msg
            loop
```

programul nu se asigura ca toate mesajele logate sunt prelucrate



➤ Exemplu: serviciu de logare

```
logMessage :: Logger -> String -> IO ()  
logMessage (Logger m) s = putMVar m s
```

```
logMessTh :: Logger -> IO()  
logMessTh log = do  
    msg <- getLine  
    if (msg == "bye")  
    then logStop log  
    else do  
        logMessage log msg  
        logMessTh log
```




-- la fel

```
initLogger :: IO Logger  
initLogger = do  
    m <- newEmptyMVar  
    let log = Logger m  
    forkIO (logger log)  
    return log  
  
main = do  
    log <- initLogger  
    logMessTh log
```



Exemplu: serviciu de logare

```
data Logger = Logger (MVar LogCommand)
data LogCommand = Message String | Stop (MVar ())
```



➤ Exemplu: serviciu de logare

```
data Logger = Logger (MVar LogCommand)
data LogCommand = Message String | Stop (MVar ())
```

```
logMessage :: Logger -> String -> IO ()
logMessage (Logger m) s = putMVar m s
```

```
logMessTh :: Logger -> IO()
logMessTh log = do
    msg <- getLine
    if (msg == "bye")
    then logStop log
    else do
        logMessage log msg
        logMessTh log
```

```
logStop :: Logger -> IO ()
logStop (Logger m) = do
    s <- newEmptyMVar
    putMVar m (Stop s)
    takeMVar s
```



Exemplu: serviciu de logare

```
logger :: Logger -> IO ()
logger (Logger m) = loop
  where loop = do
    cmd <- takeMVar m
    case cmd of
      Message msg -> do
        putStrLn ("mesaj: " ++ msg)
        loop
      Stop s -> do
        putStrLn "logger: stop"
        putMVar s ()
```

Thread-ul logger va debloca s cand cand ajunge la Stop s

```
data Logger = Logger (MVar LogCommand)
data LogCommand = Message String | Stop (MVar ())
```

```
logStop :: Logger -> IO ()
logStop (Logger m) = do
  s <- newEmptyMVar
  putMVar m (Stop s)
  takeMVar s
```

logger.hs ©2012, Simon Marlow



```
*Main> main
mes:
mes1
mesm:e
saj: mes1
mes2
memseasj::
  mes2
mes3
mesm:e
saj: mes3
bye
```

Atentie!

Accesul la stdout nu este thread-safe,
deci trebuie sincronizat

```
stdo <- newMVar ()
```

```
tswrite stdo s = do
    takeMVar stdo
    putStrLn s
    putMVar stdo ()
```

```
*Main> main
mes:
mesajul 1
mesaj: mesajul 1
mes:
mesajul 2
mes:
mesaj: mesajul 2
mesajul 3
mes:
mesaj: mesajul 3
mesajul 4
mes:
mesaj: mesajul 4
bye
```

➤ Semafoare

```
import Control.Concurrent.QSem
```

```
data QSem
```

```
newQSem :: Int -> IO Qsem
```

```
waitQSem :: QSem -> IO()    -- aquire, il ocupa
```

```
signalQSem :: QSem -> IO()  -- release, il elibereaza
```

un semafor care sincronizeaza accesul la **n** resurse se defineste astfel:

```
qs <- newQsem n
```

Exemplu: qsemrcmy.hs

O multime de taskuri acceseaza simultan o resursa reprezentata printr-un **QSem**;
pentru a se executa, fiecare task trebuie sa acceseze resursa, pe care o elibereaza la sfarsitul executiei.

```
import Control.Concurrent
import Control.Monad

main :: IO ()
main = do
    q <- newQSem 3
    stdo <- newEmptyMVar
    let workers = 5
        prints = 2 * workers
    mapM_ (forkIO . worker q m) [1..workers]
    replicateM_ prints $ takeprint stdo
```

```
takeprint :: MVar String -> IO()
```

```
takeprint stdo = do
```

```
    s <- takeMVar stdo
```

```
    print s
```

```
worker :: QSem -> MVar String -> Int -> IO ()
```

```
worker q m w= do
```

```
    waitQSem q
```

```
    putMVar stdo $ "Worker " ++ show w ++ " acquired the lock."
```

```
    threadDelay 2000000    -- microseconds
```

```
    signalQSem q
```

```
    putMVar stdo $ "Worker " ++ show w ++ "released the lock."
```

http://rosettacode.org/wiki/Metered_concurrency

q este semaforul care controleaza resursele

stdo coordoneaza accesul la stdout

<https://www.haskell.org/hoogle/>



```
Prelude> :l qsemrcmy.hs
[1 of 1] Compiling Main                ( qsemrcmy.hs, interpreted )
Ok, modules loaded: Main.
*Main> main
"Worker 1 has acquired the lock."
"Worker 2 has acquired the lock."
"Worker 3 has acquired the lock."
"Worker 2 has released the lock."
"Worker 3 has released the lock."
"Worker 1 has released the lock."
"Worker 5 has acquired the lock."
"Worker 4 has acquired the lock."
"Worker 4 has released the lock."
"Worker 5 has released the lock."
```

in Concurrent Haskell
concurenta este nedeterminista

```
*Main> main
"Worker 1 has acquired the lock."
"Worker 2 has acquired the lock."
"Worker 3 has acquired the lock."
"Worker 1 has released the lock."
"Worker 5 has acquired the lock."
"Worker 2 has released the lock."
"Worker 4 has acquired the lock."
"Worker 3 has released the lock."
"Worker 4 has released the lock."
"Worker 5 has released the lock."
```




➤ Implementarea QSem

```
type QSem = MVar (Int, [MVar ()])

newQSem :: Int -> IO QSem

newQSem n = newMVar (n,[])
            -- qsem <- newQSem 3

waitQSem :: QSem -> IO()      -- ocupa
signalQSem :: QSem -> IO()    -- elibereaza
```



n = nr. de resurse

blki = un thread care cere acces la resursa
este blocat pe variabila **blki**

daca **n > 0** atunci **qsem = (n, [])**

altfel **qsem = (0, [blk1, blk2, ...])**

Implementarea din:

Concurrent Haskell

SL Peyton Jones, A Gordon, S Finne, 1996



➤ Implementarea **QSem** - *Concurrent Haskell* SL Peyton Jones, A Gordon, S Finne, 1996

```
type QSem = MVar (Int, [MVar ()])
```

```
newQSem :: Int -> IO QSem
```

```
newQSem n = newMVar (n,[])
```

daca $n > 0$ atunci $qsem = (n, [])$
altfel $qsem = (0, [blk1, blk2, ...])$

eliberarea resursei

```
signalQSem :: QSem -> IO()
```

```
signalQSem qsem = do
```

```
    (avail,blkd) <- takeMVar qsem
```

```
    case blkd of
```

```
        [] -> putMVar qsem (avail+1,[])
```

```
        (blk:blkd') -> do
```

```
            putMVar qsem (0,blkd')
```

```
            putMVar blk ()
```

fiecare thread elibereaza variabila proprie a unui thread in asteptare

ocuparea resursei

```
waitQSem :: QSem -> IO()
```

```
waitQSem qsem = do
```

```
    (avail,blkd) <- takeMVar qsem
```

```
    if avail > 0
```

```
        then putMVar qsem (avail-1, [])
```

```
        else
```

```
            do
```

```
                blk <- newEmptyMVar
```

```
                putMVar qsem (0,blk:blkd)
```

```
                takeMVar blk – threadul e blocat pe variabila proprie
```



➤ Readers/Writers problem

- Mai multe threaduri au acces la o resursa.
- Unele threaduri scriu (writers), iar altele citesc (readers).
- Resursa poate fi accesata simultan de mai multi cititori.
- Resursa poate fi accesata de un singur scriitor.
- Resursa nu poate fi accesata simultan de cititori si de scriitori.

```
import Control.Concurrent.ReadWriteLock
new :: IO RWLock
acquireRead :: IO RWLock -> IO ()
releaseRead :: IO RWLock -> IO ()
acquireWrite :: IO RWLock -> IO ()
releaseWrite :: IO RWLock -> IO ()
```



➤ Readers/Writers problem

Mai multe threaduri au acces la o resursa.

Unele threaduri scriu (writers), iar altele citesc (readers).

Resursa poate fi accesata simultan de mai multi cititori.

Resursa poate fi accesata de un singur scriitor.

Resursa nu poate fi accesata simultan de cititori si de scriitori.

Pentru sincronizare folosim:

- un semafor binar care da acces la citit sau la scris: `writel`
- un monitor in care se inregistreaza nr. de cititori: `readL`

```
data MyRWLock = MyRWL {readL :: MVar Int, writel :: MyLock}
```

➤ Reader/Writer Lock

```
type MyLock = MVar ()  
newLock = newMVar ()  
acquireLock m = takeMVar m  
releaseLock m = putMVar m ()
```

```
data MyRWLock = MyRWL {readL :: MVar Int, writeL :: MyLock}  
newMyRWLock :: IO MyRWLock  
newMyRWLock = do  
    readL <- newMVar 0  
    writeL <- newLock  
    return (MyRWL readL writeL)
```

```
acquireWrite :: MyRWLock -> IO ()  
acquireWrite (MyRWL readL writeL) = acquireLock writeL  
  
releaseWrite :: MyRWLock -> IO ()  
releaseWrite (MyRWL readL writeL) = releaseLock writeL
```



➤ Reader/Writer Lock

```
data MyRWLock = MyRWL {readL :: MVar Int, writeL :: MyLock}
```

```
acquireRead :: MyRWLock -> IO ()
acquireRead (MyRWL readL writeL) = do
    n <- takeMVar readL -- n readers
    if (n == 0) then do
        acquireLock writeL
        putMVar readL 1
    else putMVar readL (n+1)

releaseRead :: MyRWLock -> IO ()
releaseRead (MyRWL readL writeL) = do
    n <- takeMVar readL
    if (n == 1) then do
        releaseLock writeL
        putMVar readL 0
    else putMVar readL (n-1)
```



➤ Exemplu: Readers/Writers

```
genread n rwl lib = if (n==0)
    then putStrLn "no more readers"
    else do
        reader n rwl lib
        threadDelay 20
        genread (n-1) rwl lib
genwrite n rwl lib = if (n==0)
    then putStrLn "no more writers"
    else do
        writer n rwl lib
        threadDelay 100
        genwrite (n-1) rwl lib

main = do
    lib <- newMVar 0    -- resursa
    rwl <- newMyRWLock
    forkIO $ genread 10 rwl lib
    forkIO $ genwrite 5 rwl lib
    getLine
```

```
reader i rwl lib = do
    acquireRead rwl
    c <- readMVar lib -- non blocking
    putStrLn $ (show i) ++ (show c)
    -- "Reader " ++ (show i) ++ " reads: " ++ (show c)
    releaseRead rwl

writer i rwl lib = do
    acquireWrite rwl
    putStrLn $ show i
    -- "Writer " ++ (show i) ++ " writes " (show i)
    c <- takeMVar lib
    putMVar lib i
    releaseWrite rwl
```



➤ Readers/Writers

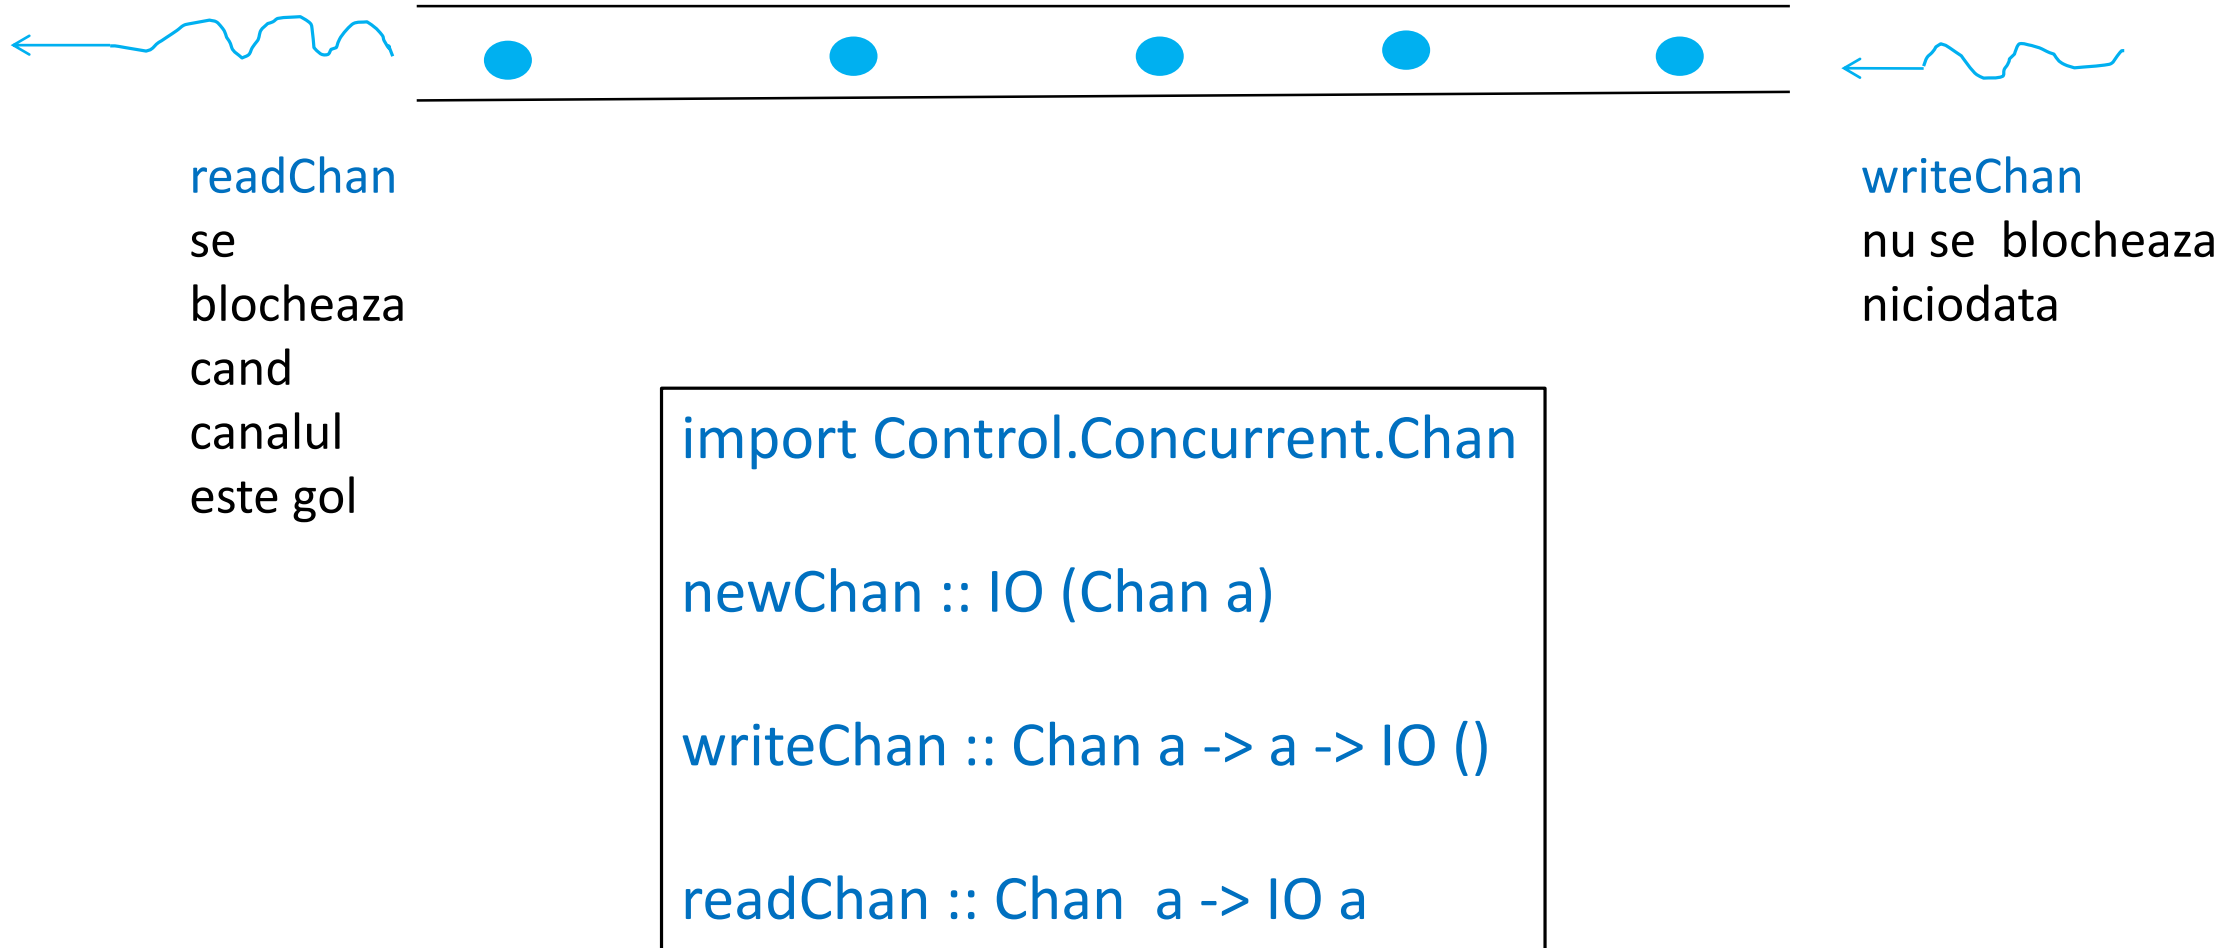
```
genread n rwl lib = if (n==0)
    then putStrLn "no more readers"
    else do
        reader n rwl lib
        threadDelay 20
        genread (n-1) rwl lib
genwrite n rwl lib = if (n==0)
    then putStrLn "no more writers"
    else do
        writer n rwl lib
        threadDelay 100
        genwrite (n-1) rwl lib

main = do
    lib <- newMVar 0    -- resursa
    rwl <- newMyRWLock
    forkIO $ genread 10 rwl lib
    forkIO $ genwrite 5 rwl lib
    getLine
```

```
Prelude> :l myrw.hs
[1 of 1] Compiling Main
Ok, modules loaded: Main.
*Main> main
Reader 10 reads: 0
Writer 5 writes 5
Reader 9 reads: 5
Reader 8 reads: 5
Writer 4 writes 4
Reader 7 reads: 4
Reader 6 reads: 4
Writer 3 writes 3
Reader 5 reads: 3
Writer 2 writes 2
Reader 4 reads: 2
Writer 1 writes 1
Reader 3 reads: 1
no more writers
Reader 2 reads: 1
Reader 1 reads: 1
no more readers
```



➤ Canale de comunicare: canale implementate cu MVar



➤ Exemplu: doua canale: **cin** si **cout**

- thread –ul parinte citeste siruri si le pune pe canalul **cin**.
- un thread citeste sirurile de pe **cin**, le imparte in cuvinte iar cuvintele le pune pe canalul **cout**.
- un alt thread ia cuvintele de pe **cout**, si le scrie la iesire cu litere mari

```
import Control.Monad
import Control.Concurrent
import Data.Char

mymain = do
    cin <- newChan
    cout <- newChan
    forkIO $ forever (move cin cout)
    forkIO $ forever (upout cout)
    load cin
```

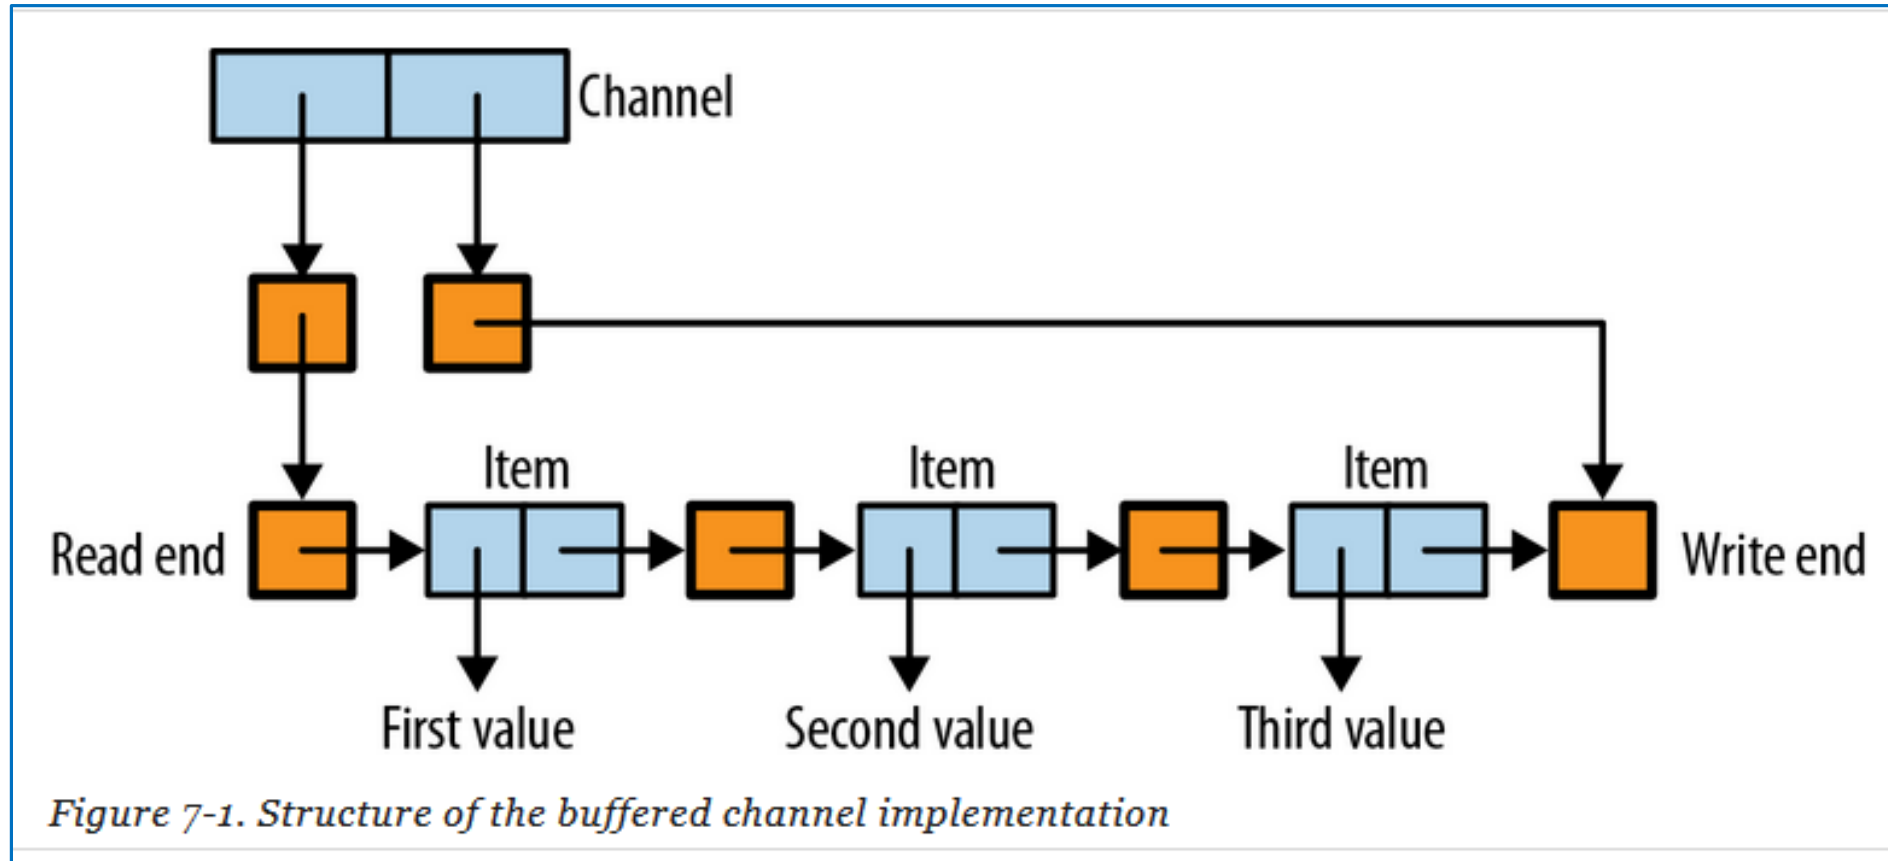
```
move c1 c2 = do
    v1 <- readChan c1
    let ls = words v1
    mapM_ (writeChan c2) ls

upout c = do
    str <- readChan c
    putStrLn (map toUpper str)

load c = do
    str <- getLine
    if (str == "exit")
    then return()
    else do
        writeChan c str
        load
```



➤ Canale de comunicare formate din variabile MVar



http://chimera.labs.oreilly.com/books/1230000000929/ch07.html#sec_channels



➤ Canale formate din variabile MVar

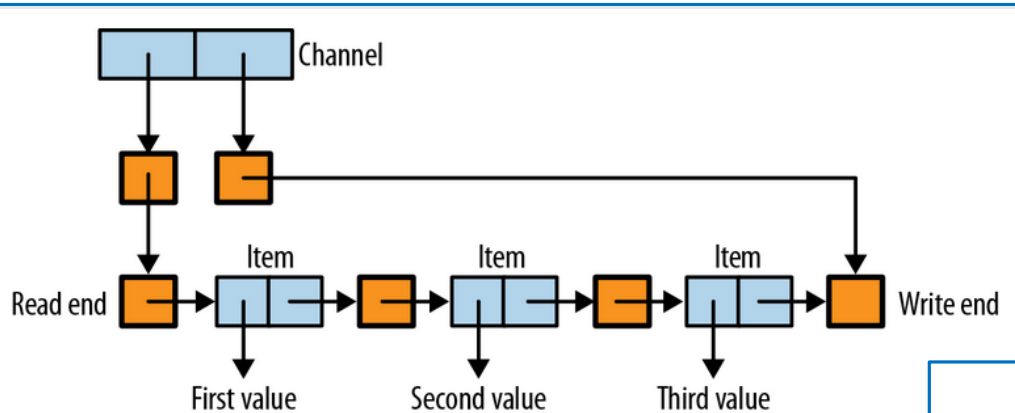



Figure 7-1. Structure of the buffered channel implementation

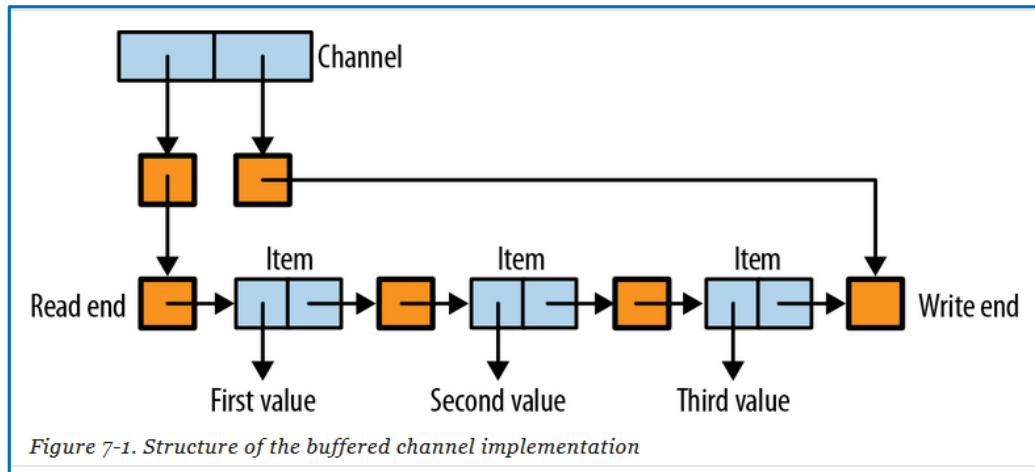
```
type Stream a = MVar (Item a)
data Item a   = Item a (Stream a)
```



```
data Chan a = Chan (MVar (Stream a)) (MVar (Stream a))
```

```
c <- newChan
v <- readChan c
putChan c v
```

chan.hs ©2012, Simon Marlow



```

type Stream a = MVar (Item a)
data Item a   = Item a (Stream a)
data Chan a   = Chan (MVar (Stream a)) (MVar (Stream a))

c <- newChan
v <- readChan c
putChan c v


```

"If multiple threads concurrently call `readChan`, the first one will successfully call `takeMVar` on the read end, but the subsequent threads will all block at this point until the first thread completes the operation and updates the read end. If multiple threads call `writeChan`, a similar thing happens: the write end of the `Chan` is the synchronization point, allowing only one thread at a time to add an item to the channel. However, the read and write ends, being separate `MVars`, allow concurrent `readChan` and `writeChan` operations to proceed without interference."

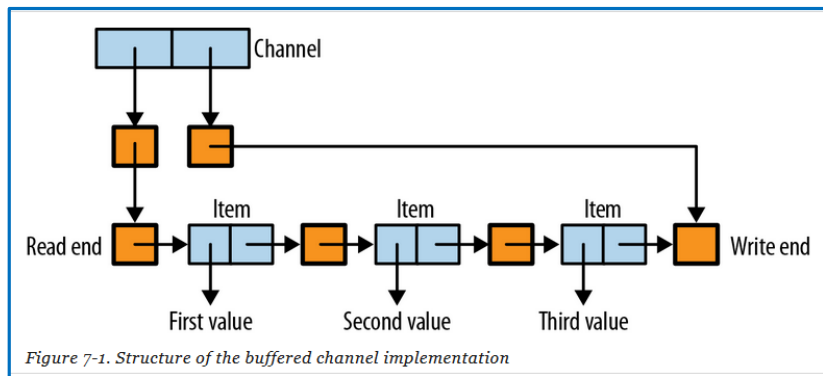
http://chimera.labs.oreilly.com/books/1230000000929/ch07.html#sec_channels

<https://www.haskell.org/hoogle/>





➤ Implementarea canalelor



```
newChan :: IO(Chan a)
```

```
newChan = do
```

```
    emptyStream <- newEmptyMVar
    readVar <- newMVar emptyStream
    writeVar <- newMVar emptyStream
    return (Chan readVar writeVar)
```

contine **Item**-ul care
va fi citit

contine variabila in care
se va scrie noul **Item**

```
type Stream a = MVar (Item a)
data Item a = Item a (Stream a)
data Chan a = Chan (MVar (Stream a)) (MVar (Stream a))
```

```
readChan :: Chan a -> IO a
```

```
readChan (Chan rV wV) = do
```

```
    stream <- takeMVar rV
```

```
    Item val str <- takeMVar stream
```

```
    putMVar rV <- str
```

```
    return val
```

```
writeChan :: Chan a -> a -> IO()
```

```
writeChan (Chan rV wV) val = do
```

```
    newStream <- newEmptyMVar
```

```
    writeEnd <- takeMVar wV
```

```
    putMVar writeEnd (Item val newStream)
```

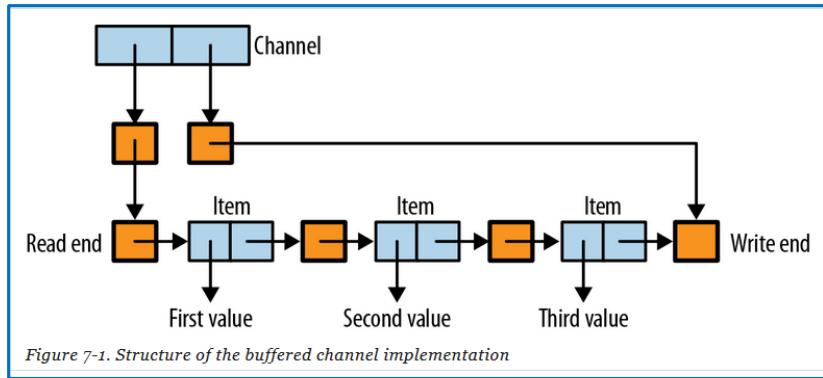
```
    putMVar wV newStream
```

http://chimera.labs.oreilly.com/books/1230000000929/ch07.html#sec_channels

<https://www.haskell.org/hoogle/>



➤ Implementarea canalelor



```
type Stream a = MVar (Item a)
data Item a = Item a (Stream a)
data Chan a = Chan (MVar (Stream a)) (MVar (Stream a))
newChan :: IO (Chan a)
writeChan :: Chan a -> a -> IO ()
readChan :: Chan a -> IO a
```

`dupChan :: Chan a -> IO (Chan a)`

- noul canal este initial gol
- dupa crearea canalului duplicat, ceea ce se scrie pe oricare din canale poate fi citit de pe oricare cele doua canale
- citirea de pe un canal nu elimina elementul de pe celalalt canal.

```
main = do c <- newChan
          writeChan c 'a'
          readChan c >=> print
          c2 <- dupChan c
          writeChan c 'b'
          readChan c >=> print
          readChan c2 >=> print
```

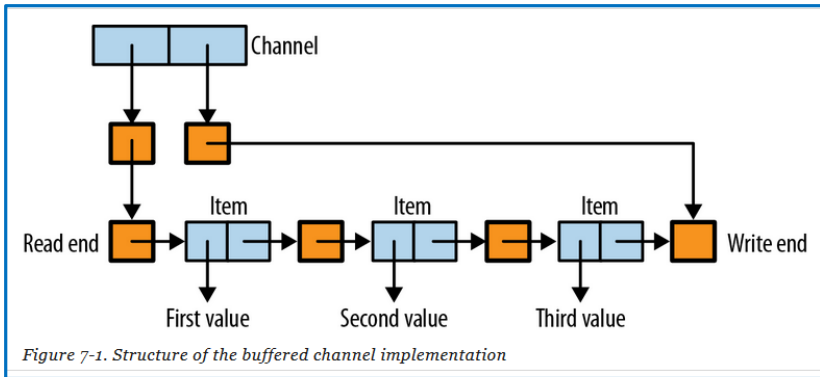
```
Prelude> :l chan2.hs
[1 of 1] Compiling Main
Ok, modules loaded: Main.
*Main> main
'a'
'b'
'b'
```

http://chimera.labs.oreilly.com/books/1230000000929/ch07.html#sec_channels

<https://www.haskell.org/hoogle/>



➤ Implementarea canalelor



```
type Stream a = MVar (Item a)
data Item a = Item a (Stream a)
data Chan a = Chan (MVar (Stream a)) (MVar (Stream a))
newChan :: IO (Chan a)
writeChan :: Chan a -> a -> IO ()
readChan :: Chan a -> IO a
```

```
dupChan :: Chan a -> IO (Chan a)
dupChan (Chan _ writeVar) = do
    writeEnd <- readMVar writeVar
    newReadVar <- newMVar writeEnd
    return (Chan newReadVar writeVar)
```

```
readChan :: Chan a -> IO a
readChan (Chan rV wV) = do
    stream <- takeMVar rV
    Item val str <- readMVar stream --!!
    putMVar rV str
    return val
```

readMVar este necesar deoarece continutul trebuie sa ramana accesibil celuilalt canal.

```
readMVar :: MVar a -> IO a
readMVar m = do
    v <- takeMVar m
    putMVar m v
    return v
```

http://chimera.labs.oreilly.com/books/1230000000929/ch07.html#sec_channels

<https://www.haskell.org/hoogle/>



“**Concurrent computing** is a form of computing in which several computations are executing during overlapping time periods—*concurrently*—instead of *sequentially* (one completing before the next starts)[...]

A concurrent system is one where a computation can advance without waiting for all other computations to complete; where more than one computation can advance at the same time.”

Operating System Concepts 9th edition, Abraham Silberschatz

Exemplu: incarcarea mai multor pagini web

secvential



concurent



```
import Data.ByteString as B
import GetURL    -- parconc-examples
main = do
    m1 <- newEmptyMVar
    forkIO $ do
        r <- getURL "http://..."
        putMVar m1 r

    m2 <- newEmptyMVar
    forkIO $ do
        r <- getURL "http://..."
        putMVar m2 r

    r1 <- takeMVar m1
    r2 <- takeMVar m2

    print (B.length r1, B.length r2)
```



➤ Comunicare asincrona

Se creaza un thread separat pentru fiecare actiune si se asteapta rezultatul

```
m1 <- newEmptyMVar
forkIO $ do
    r <- getURL "http://www.fmi.ro "
    putMVar m1 r
r1 <- takeMVar m
```

```
a <- async (getURL "http://www.fmi.ro ")
r <- wait a
```

```
data Async a = Async (MVar a)
```

```
async :: IO a -> IO (Async a)
```

```
async action = do
```

```
    var <- newEmptyMVar
```

```
    forkIO (do r <- action; putMVar var r)
```

```
    return (Async var)
```

```
wait :: Async a -> IO a
```

```
wait (Async var) = readMVar var
```

readMVar nu blocheaza threadul, deci mai multe apeluri **wait** pot fi facute pentru aceeasi operatie asincrona



➤ Comunicare asincrona

```
import Control.Concurrent
import Text.Printf
import qualified Data.ByteString as B
import GetURL -- parconc-examples
import Timelt  -- parconc-examples
```

```
timeDownload :: String -> IO ()
timeDownload url = do
    (page, time) <- timeit $ getURL url
    printf " %s (%d bytes, %.2fs)\n" url (B.length page) time
```

```
sites = ["url1", "url2", ...]
main = do
    as <- mapM (async . timeDownload) sites
    mapM_ wait as
```

```
data Async a = Async (MVar a)

async :: IO a -> IO (Async a)
async action = do
    var <- newEmptyMVar
    forkIO (do r <- action; putMVar var r)
    return (Async var)

wait :: Async a -> IO a
wait (Async var) = readMVar var
```

geturl3.hs ©2012, Simon Marlow



➤ Comunicare asincrona

In exemplul anterior vrem sa scriem un mesaj cand s-a descarcat prima pagina

```
sites = ["url1", "url2", ...]
```

```
download m url = do
```

```
    r <- getURL url
```

```
    putMVar m (url, r)
```

```
main :: IO ()main = do
```

```
    m <- newEmptyMVar
```

```
    mapM_ (forkIO . download m) sites
```

```
    (url, r) <- takeMVar m
```

```
    printf "%s was first (%d bytes)\n" url (B.length r)
```

```
    replicateM_ (length sites - 1) (takeMVar m)
```

threadul principal va accesa
variabila `m` in momentul
in care primeste o valoare



➤ Async - comunicare asincrona (folosind MVar)

```
import Control.Concurrent
import Text.Printf
import qualified Data.ByteString as B
import GetURL -- parconc-examples
import TimeIt -- parconc-examples
```

```
timeDownload :: String -> IO ()
timeDownload url = do
    (page, time) <- timeit $ getURL url
    printf " %s (%d bytes, %.2fs)\n" url (B.length page) time
```

```
data Async a = Async (MVar a)

async :: IO a -> IO (Async a)
async action = do
    var <- newEmptyMVar
    forkIO (do r <- action; putMVar var r)
    return (Async var)
```

```
wait :: Async a -> IO a
wait (Async var) = readMVar var
```

```
main = do
    as <- mapM (async . timeDownload) sites -- sites = ["url1", "url2", ...]
    mapM_ wait as
```

asteapta ca toate actiunile asincrone sa se termine, monitorizand fiecare actiune in parte; un alt thread ar putea interveni inainte ca toate actiunile sa se termine

Vom rezolva acesta problema folosind STM

