
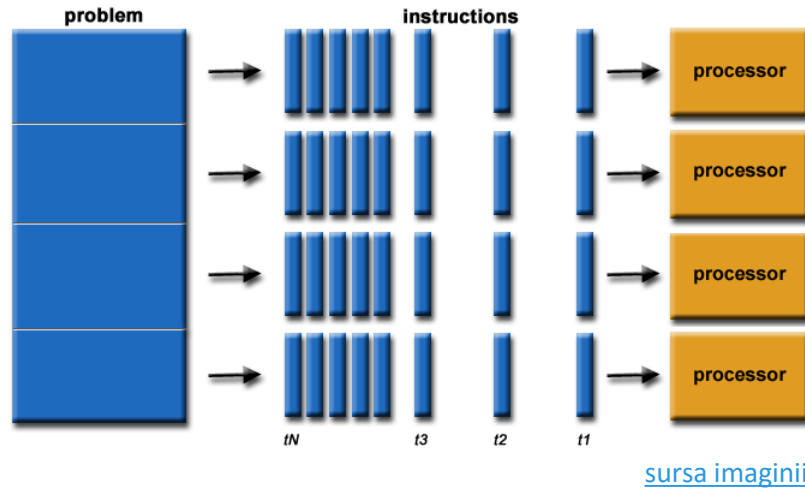


IMPLEMENTAREA CONCURENTEI IN LIMBAJE DE PROGRAMARE

Parallelism in Haskell introducere

Ioana Leustean



[Chapter 2: Basic Parallelism](#)

Parallelism is about speeding up a program by using multiple processors.

Pure Parallelism (Control.Parallel): Speeding up a pure computation using multiple processors.

Pure parallelism has these advantages:

- Guaranteed deterministic (same result every time)
- no race conditions or deadlocks.

Concurrency (Control.Concurrent): multiple threads of control that execute "at the same time".

- Threads are in the IO monad.
- IO operations from multiple threads are interleaved non-deterministically
- communication between threads must be explicitly programmed.
- Threads may execute on multiple processors simultaneously.
- Danger: race conditions and deadlocks.

<https://wiki.haskell.org/Parallelism>

<https://www.haskell.org/hoogle/>



- Paralelismul in Haskell se obtine folosind monada `Eval` care este definite in modulul `Control.Parallel.Strategies`

```
data Eval a
instance Monad Eval
```

```
runEval :: Eval a -> a
```

argumentul poate fi
evaluat in paralel

```
rpar :: a -> Eval a
```

evalueaza argumentul si
asteapta rezultatul

```
rseq :: a -> Eval a
```

```
import Control.Parallel.Strategies
```

```
-- fib n  Fibonacci
```

```
test = runEval $ do
```

```
    x <- rpar (fib 36)
```

```
    y <- rpar (fib 37)
```

```
    z <- rseq (fib 33)
```

```
    return (x,y,z)
```

```
main = print test
```



```
*Main> :! ghc --make mypar1.hs
[1 of 1] Compiling Main                ( mypar1.hs, mypar1.o )
Linking mypar1.exe ...
```

mypar1.hs

```
import Control.Parallel.Strategies
```

```
-- fib n Fibonacci
```

```
test = runEval $ do
    x <- rpar (fib 36)
    y <- rpar (fib 37)
    z <- rseq (fib 33)
    return (x,y,z)
```

```
main = print test
```

```
D:\DIR\HS\myexc>mypar1 +RTS -s
(24157817,39088169,5702887)
 17,229,615,344 bytes allocated in the heap
 21,906,856 bytes copied during GC
 60,008 bytes maximum residency <4 sample(s)>
 22,080 bytes maximum slop
 2 MB total memory in use <0 MB lost due to fragmentation>

Gen  0      32936 colls,    0 par    0.11s   0.15s   0.0000s   0.0009s
Gen  1         4 colls,    0 par    0.00s   0.00s   0.0001s   0.0002s

INIT      time    0.00s   < 0.00s elapsed>
MUT       time    8.24s   < 8.28s elapsed>
GC         time    0.11s   < 0.15s elapsed>
EXIT      time    0.00s   < 0.00s elapsed>
Total     time    8.36s   < 8.43s elapsed>

%GC        time      1.3%   <1.7% elapsed>

Alloc rate 2,091,771,670 bytes per MUT second

Productivity 98.7% of total user, 97.9% of total
```

```
> mypar1 +RTS -s
```

CPU time: 8.36s/ wall-clock time: 8.43s

Total time 8.36s (8.43s elapsed)



```
*Main> :! ghc --make -threaded mypar1t.hs
[1 of 1] Compiling Main             ( mypar1t.hs, mypar1t.o )
Linking mypar1t.exe ...
```

mypar1t.hs

```
import Control.Parallel.Strategies
```

```
-- fib n Fibonacci
```

```
test = runEval $ do
    x <- rpar (fib 36)
    y <- rpar (fib 37)
    z <- rseq (fib 33)
    return (x,y,z)
```

```
main = print test
```

```
D:\DIR\HS\myexc>mypar1t +RTS -N2 -s
<24157817,39088169,5702887>
 17,229,689,904 bytes allocated in the heap
 22,353,560 bytes copied during GC
 128,216 bytes maximum residency <4 sample(s)>
 30,296 bytes maximum slop
 3 MB total memory in use <0 MB lost due to fragmentation>

Gen  0      21779 colls, 21779 par  0.17s  0.14s  0.0000s  0.0093s
Gen  1         4 colls,   3 par  0.00s  0.00s  0.0001s  0.0002s

Parallel GC work balance: 42.83% <serial 0%, perfect 100%>
TASKS: 4 <1 bound, 3 peak workers <3 total>, using -N2>
SPARKS: 2 <2 converted, 0 overflowed, 0 dud, 0 GC'd, 0 fizzled>

INIT    time    0.00s  < 0.00s elapsed>
MUT     time    9.28s  < 5.87s elapsed>
GC       time    0.17s  < 0.14s elapsed>
EXIT    time    0.00s  < 0.00s elapsed>
Total   time    9.45s  < 6.01s elapsed>
```

```
> mypar1 +RTS -N2 -s
```

```
-----
SPARKS: 2 (2 converted,...)
Total time 9.45 s (6.01s elapsed)
```

bytes per MUT second

converted = folosite efectiv in paralelism

speedup 8.43/6.01 = 1.4



➤ Sparks

"The argument to `rpar` is called a spark. The runtime collects sparks in a pool and uses this as a source of work when there are spare processors available, using a technique called work stealing. Sparks may be evaluated at some point in the future, or they might not—it all depends on whether there is a spare core available. Sparks are very cheap to create: `rpar` essentially just writes a pointer to the expression into an array."

http://chimera.labs.oreilly.com/books/1230000000929/ch02.html#sec_par-eval-sudoku2

"Sparks are specific to parallel Haskell. Abstractly, a spark is a pure computation which may be evaluated in parallel. Whether or not a spark is evaluated in parallel with other computations, or other Haskell IO threads, depends on what your hardware supports and on how your program is written. Sparks are put in a work queue and when a CPU core is idle, it can execute a spark by taking one from the work queue and evaluating it.

On a multi-core machine, both threads and sparks can be used to achieve parallelism. Threads give you concurrent, non-deterministic parallelism, while sparks give you **pure deterministic parallelism**. Haskell threads are ideal for applications like network servers where you need to do lots of I/O and using concurrency fits the nature of the problem. Sparks are ideal for speeding up pure calculations where adding non-deterministic concurrency would just make things more complicated. "

<https://wiki.haskell.org/Parallel/Glossary>



Here are some other things that can happen to a spark:

converted

Are turned into real parallelism at runtime

dud

When `rpar` is applied to an expression that is already evaluated, this is counted as a dud and the `rpar` is ignored.

GC'd

The sparked expression was found to be unused by the program, so the runtime removed the spark. We'll discuss this in more detail in "GC'd Sparks and Speculative Parallelism".

fizzled

The expression was unevaluated at the time it was sparked but was later evaluated independently by the program. Fizzled sparks are removed from the spark pool.

"It is possible that a spark in the spark pool can refer to a computation that has already been evaluated by the program. Perhaps there were not enough processors to evaluate the spark in parallel, and another thread ended up evaluating the computation during the normal course of computing its results.

When a spark in the spark pool refers to a value, rather than an unevaluated computation, we say the spark has fizzled; this potential for parallel execution has expired"

Seq no more: Better Strategies for Parallel Haskell, 2010

<http://research.microsoft.com/pubs/138042/haskell18-marlow.pdf>

http://chimera.labs.oreilly.com/books/1230000000929/ch02.html#sec_par-eval-sudoku2

<https://www.haskell.org/hoogle/>



```
import Control.Parallel.Strategies
import System.Environment
```

```
fib :: Integer -> Integer
fib 0 = 1
fib 1 = 1
fib n = fib (n-1) + fib (n-2)
```

```
test n1 n2 n3 = runEval $ do
  x <- rpar (fib n1)
  y <- rpar (fib n2)
  z <- rseq (fib n3)
  return (x,y,z)
```

```
main = do
  [s1,s2,s3] <- getArgs
  let n1 = (read s1)::Integer
      n2 = (read s2)::Integer
      n3 = (read s3)::Integer
  print (test n1 n2 n3)
```

```
Prelude> :m + System.Environment
Prelude System.Environment> :t getArgs
getArgs :: IO [String]
```

```
*Main> :! ghc --make -threaded mypar2t.hs
[1 of 1] Compiling Main                ( mypar2t.hs,
Linking mypar2t.exe ...
```

```
D:\DIR\HS\myexc>mypar2t 25 30 10 +RTS -N2 -s
<121393,1346269,89>
 366,926,856 bytes allocated in the heap
 476,568 bytes copied during GC
 94,792 bytes maximum residency <2 sample(s)>
 26,192 bytes maximum slop
 2 MB total memory in use <0 MB lost due to fragmentation>

Gen  0      642 colls,    642 par    0.03s   0.03s   0.0000s
Gen  1       2 colls,     1 par    0.00s   0.00s   0.0001s

Parallel GC work balance: 8.43% <serial 0%, perfect 100%>
TASKS: 4 <1 bound, 3 peak workers <3 total>, using -N2>
SPARKS: 2 <1 converted, 0 overflowed, 0 dud, 0 GC'd, 1 fizzled>

INIT    time    0.00s   < 0.00s elapsed>
MUT     time    0.16s   < 0.16s elapsed>
GC      time    0.03s   < 0.03s elapsed>
EXIT    time    0.00s   < 0.00s elapsed>
Total   time    0.20s   < 0.19s elapsed>
```

```
> mypar2t 25 30 10 +RTS -N2 -s
```



➤ mai mult paralelism

```
myparMap :: (a -> b) -> [a] -> Eval [b]
```

```
myparMap f [] = return []  
myparMap f (a:as) = do  
    b <- rpar (f a)  
    bs <- myparMap f as  
    return (b:bs)
```

```
import Control.Parallel.Strategies  
--import System.Environment
```

```
arg = [34,35,25, 27]
```

```
fib :: Integer -> Integer  
fib 0 = 1  
fib 1 = 1  
fib n = fib (n-1) + fib (n-2)
```

```
myparMap :: (a -> b) -> [a] -> Eval [b]  
myparMap f [] = return []  
myparMap f (a:as) = do  
    b <- rpar (f a)  
    bs <- myparMap f as  
    return (b:bs)
```

```
test = runEval (myparMap fib arg)
```

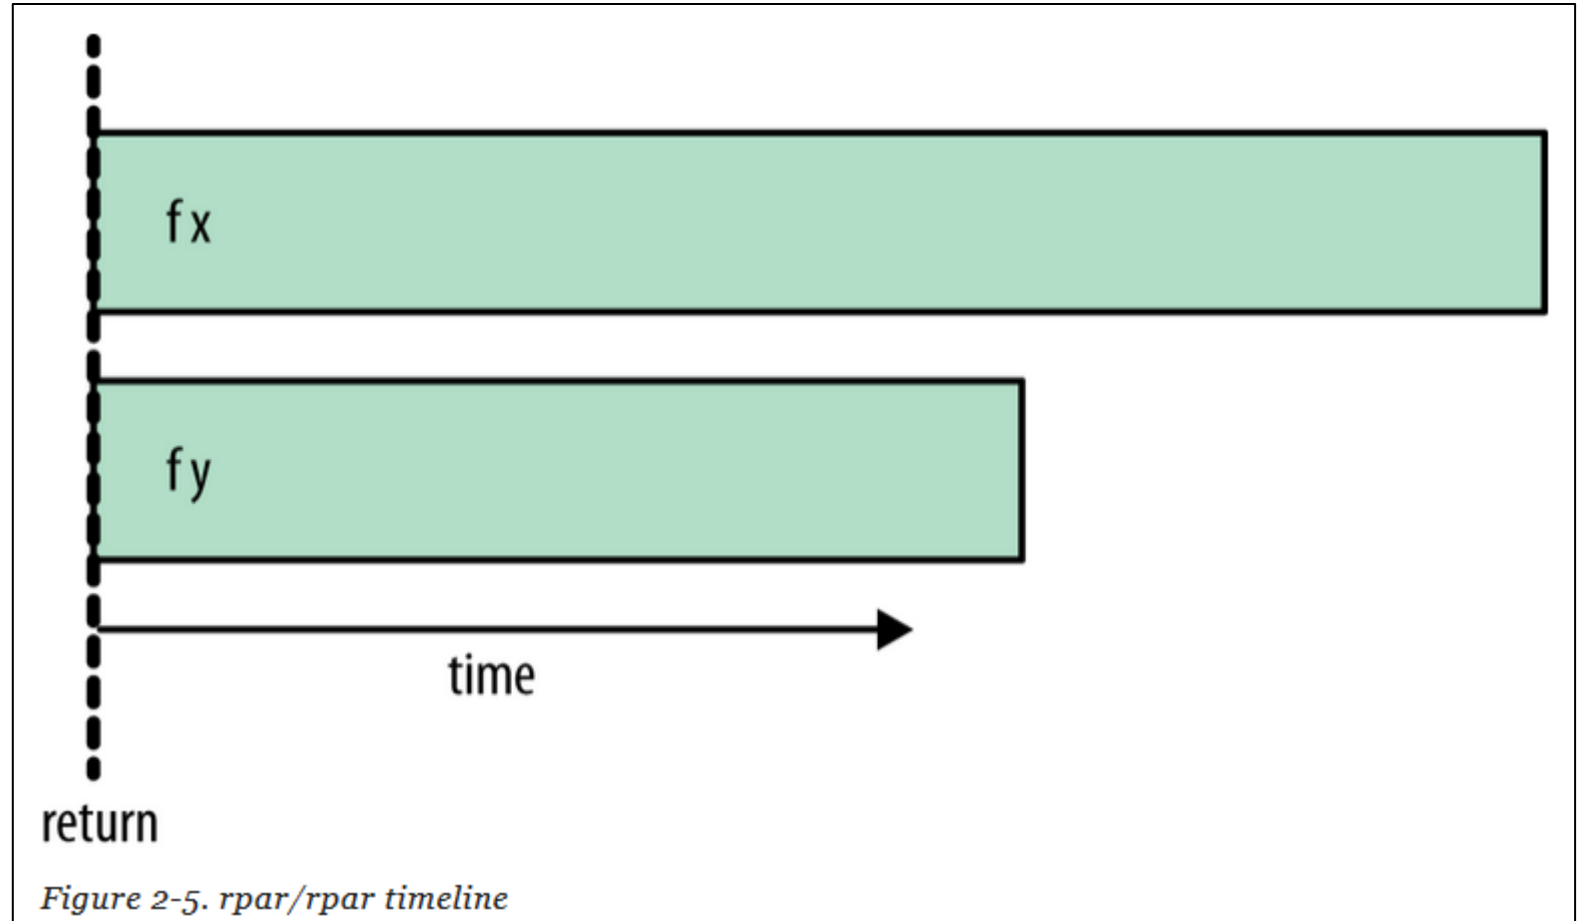
```
main = print test
```



➤ `rpar` `si` `rseq`

`rpar` "my argument can be created in parallel"
`rseq` "evaluate my argument and wait for the result"

```
test = runEval $ do
  x <- rpar (fib 36)
  y <- rpar (fib 37)
  return (x,y)
```



http://chimera.labs.oreilly.com/books/1230000000929/ch02.html#sec_par-rpar-rseq

<https://www.haskell.org/hoogle/>

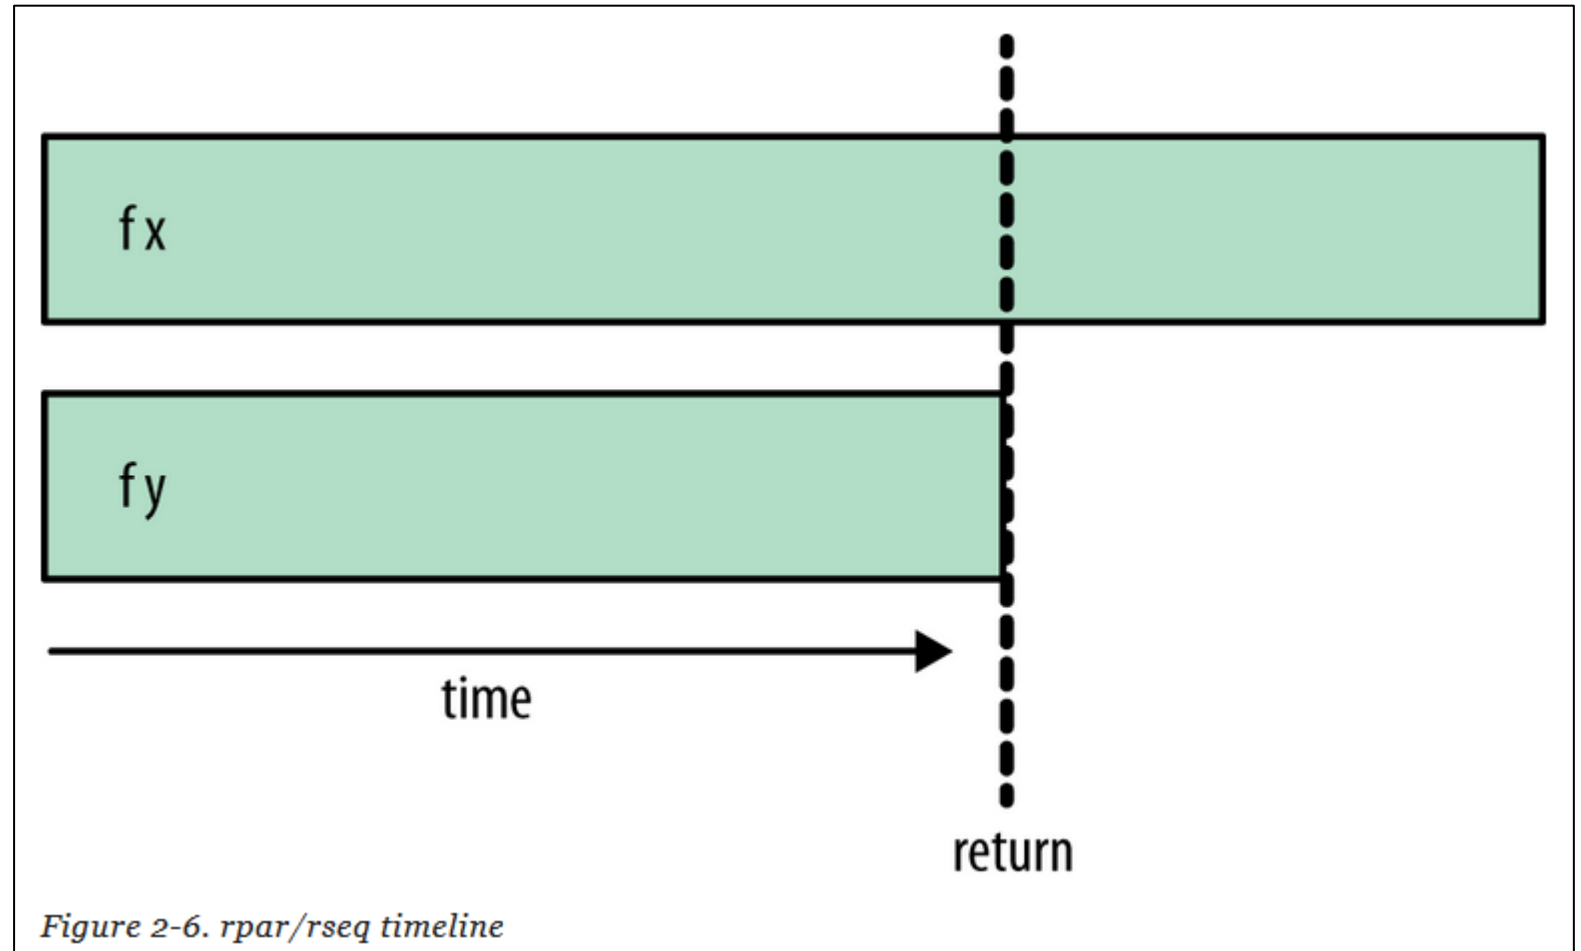


➤ rpar si rseq

rpar "my argument can be created in parallel"
rseq "evaluate my argument and wait for the result"

```
test = runEval $ do
  x <- rpar (fib 37)
  y <- rseq (fib 36)
  return (x,y)
```

nu e indicata, deoarece
nu stim care din operatii
se termina prima!



http://chimera.labs.oreilly.com/books/12300000000929/ch02.html#sec_par-rpar-rseq

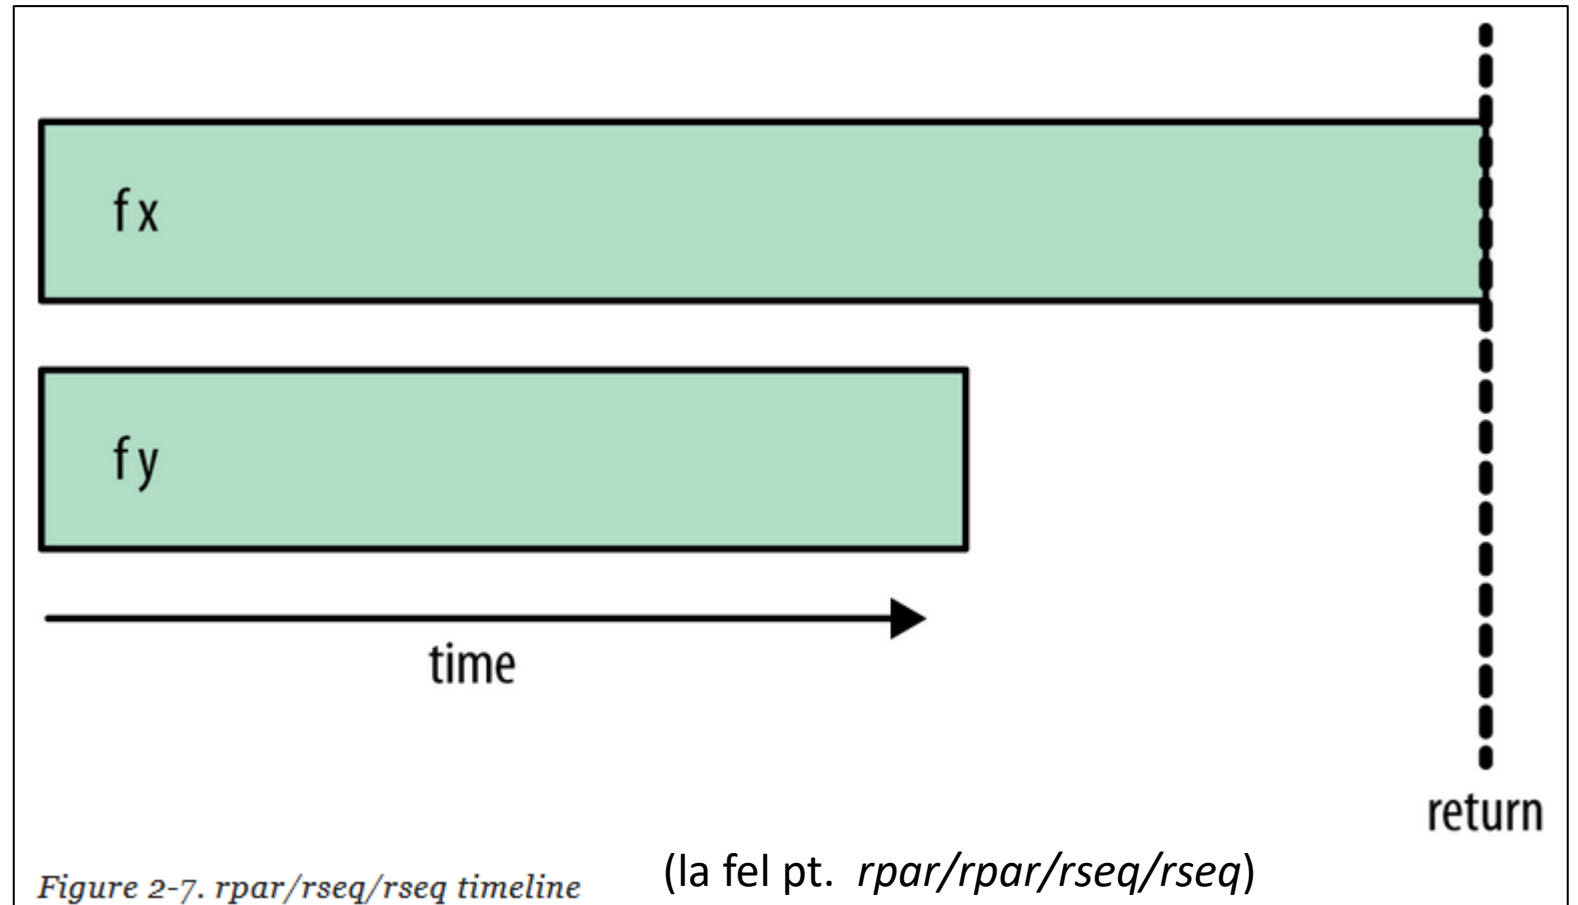
<https://www.haskell.org/hoogle/>



➤ rpar si rseq

rpar "my argument can be created in parallel"
rseq " evaluate my argument and wait for the result"

```
test = runEval $ do
  x <- rpar (fib 37)
  y <- rpar (fib 36)
  rseq x
  rseq y
  return (x,y)
```



http://chimera.labs.oreilly.com/books/1230000000929/ch02.html#sec_par-rpar-rseq



- Thunk (o expresie neevaluata)
- WHNF (weak head normal form)

```
Prelude> let x = 1+2 :: Int
```

```
Prelude> :sprint x
```

```
x = _
```

```
Prelude> let y =(x,x)
```

```
Prelude> :sprint y
```

```
y = _
```

```
Prelude> :m + Control.Exception
```

```
Prelude Control.Exception> ye <- evaluate (x,x)
```

```
Prelude Control.Exception> :sprint ye
```

```
ye = (_,_)
```

```
Prelude Control.Exception> xe <-evaluate(x)
```


```
Prelude Control.Exception> :sprint xe
```

```
xe = 3
```


`:sprint` scrie valoarea unei expresii fara a-l forta evaluarea

```
Prelude Control.Exception> :t evaluate
evaluate :: a -> IO a
```


`evaluate` isi evalueaza argumentul la whnf atunci cand actiunea IO este executata



x este un "thunk"




y este o "whnf"



z este o forma normala

➤ Thunk. WHNF

- Un thunk este o expresie neevaluata
- O expresie este in whnf daca, in reprezentarea grafica, radacina arborelui este un constructor
 - + 1 2 nu este whnf
 - 3: x:[] este whnf
- O expresie este in forma normal daca este complet evaluata.



http://en.wikibooks.org/wiki/Haskell/Laziness#Thunks_and_Weak_head_normal_form



➤ WHNF

`rpar` si `rseq` isi evalueaza argumentele la whnf

```
Prelude Control.Exception Control.Parallel.Strategies> yer <- evaluate $ runEval (rpar (x,x))
Prelude Control.Exception Control.Parallel.Strategies> :sprint yer
yer = (_,_)
Prelude Control.Exception Control.Parallel.Strategies> yeq <- evaluate $ runEval (rseq (x,x))
Prelude Control.Exception Control.Parallel.Strategies> :sprint yeq
yeq = (_,_)
```

```
Prelude Control.Exception Control.Parallel.Strategies> let ls = map (+1)[1,2,3]
Prelude Control.Exception Control.Parallel.Strategies> :sprint ls
ls = _
Prelude Control.Exception Control.Parallel.Strategies> lse <- evaluate $ runEval (rpar ls)
Prelude Control.Exception Control.Parallel.Strategies> :sprint lse
lse = _ : _
```



➤ force

`rpar` si `rseq` isi evalueaza argumentele la `whnf`

```
Prelude Control.Exception Control.Parallel.Strategies> let ls = map (+1)[1,2,3]
Prelude Control.Exception Control.Parallel.Strategies> :sprint ls
ls = _
Prelude Control.Exception Control.Parallel.Strategies> lse <- evaluate $ runEval (rpar ls)
Prelude Control.Exception Control.Parallel.Strategies> :sprint lse
lse = _ : _
```

Daca dorim ca `ls` sa fie complet evaluate de `rpar` atunci putem sa "fortam" evaluarea folosind functia `force`

```
import Control.DeepSeq
force :: NFData a => a -> a
```

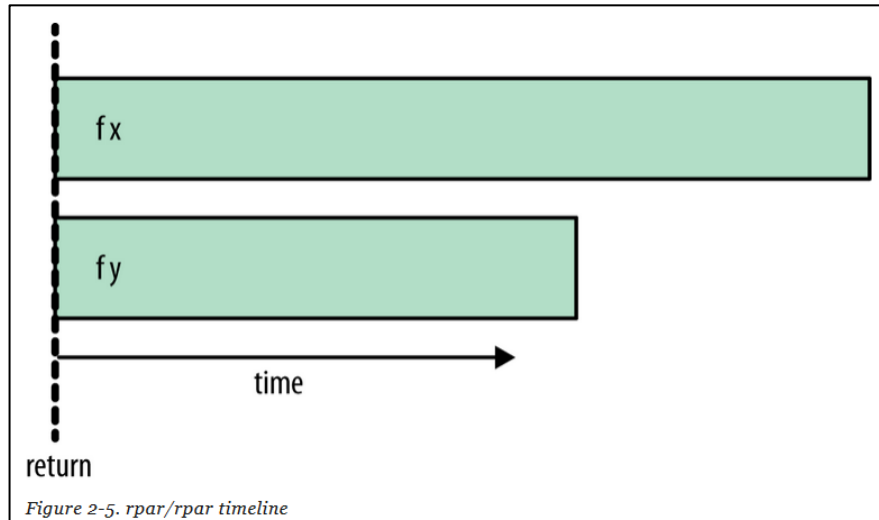
```
> lse1 <- evaluate $ runEval (rpar (force ls))
> :sprint lse1
> [2,3,4]
```

```
Prelude Control.Exception Control.Parallel.Strategies> :m + Control.DeepSeq
Prelude Control.Exception Control.Parallel.Strategies Control.DeepSeq> lse1 <- evaluate $ runEval (rpar (force ls))
Prelude Control.Exception Control.Parallel.Strategies Control.DeepSeq> :sprint lse1
lse1 = [2,3,4]
Prelude Control.Exception Control.Parallel.Strategies Control.DeepSeq> :t force
force :: NFData a => a -> a
```



➤ Strategii de evaluare

`rpar/rpar` poate fi vazuta
ca o strategie de evaluare



```
*Main> :t rpar  
rpar :: a -> Eval a
```

```
type Strategy a = a -> Eval a
```

```
parPair :: Strategy (a,b)
```

```
parPair (a,b) = do
```

```
    a' <- rpar a
```

```
    b' <- rpar b
```

```
    return (a',b')
```

```
main = print pair
```

```
    where
```

```
        pair = runEval ( parPair (fib 35, fib 36))
```

Strategiile sunt impementate in modulul [Control.Parallel.Strategies](https://hackage.haskell.org/package/Control.Parallel.Strategies)



Pentru a ilustra posibilitatile de implementare si extindere a functionalitatilor vom crea operatii noi folosind din modulul [Control.Parallel.Strategies](#) numai `rpar`

```
type MyStrategy a = a -> Eval a
```

```
myrparforce :: NFData a => MyStrategy a  
myrparforce x = rpar (force x)
```

o strategie pentru evaluare in paralel,
argumentul fiind complet evaluat



Strategii parametrizate

```
myevalPair :: MyStrategy a -> MyStrategy b -> MyStrategy (a,b)
myevalPair st1 st2 (x,y) = do
    x' <- st1 x
    y' <- st2 y
    return (x',y')
```

st1 este argument de tip **MyStrategy a**
st2 este argument de tip **MyStrategy b**

```
myparPair :: (NFData a, NFData b) => MyStrategy (a,b)
myparPair = myevalPair myrparforce myrparforce
```

st1 = myrparforce a
st2 = myrparforce b

```
main = print pair
      where
        pair = runEval ( myparPair (fib 36, fib 35))
```



Strategii parametrizate

```
myparList :: NFData a => MyStrategy a -> MyStrategy [a]
myparList strat [] = return []
myparList strat (x:xs) = do
    x' <- strat x
    xs' <- myparList strat xs
    return (x':xs')
```

strat este un parametru de tip **MyStrategy a**

myparList este o strategie de evaluare a listelor,
parametrizata de strategia de evaluare a elementelor

